

# Heath Reinhard

AI Evals & Knowledge Quality · Rubric Design & Calibration · Closed-Loop Answer Quality Systems  
South Lake Tahoe, CA · 718.541.3367 · heath.reinhard@gmail.com

[LinkedIn](#) · [Portfolio](#)

## SUMMARY

---

Documentation engineer specializing in AI evals and knowledge quality, with 7+ years in technical documentation, remote preferred. Over 4+ years fully remote at Meta, I hand-graded the verdict-balanced calibration set that anchored a production LLM judge — scoring 200 prompts across five rubric dimensions against a four-verdict vocabulary — then built the evaluation pipeline that scaled that standard, grading production AI agent outputs, classifying coverage failures, and routing corrected next steps through review and back into the knowledge base. That flywheel lifted eval pass rates 9.3pp across supported products. I supported six product teams spanning 150+ internal products as their documentation engineer, authoring the runbooks and governance frameworks their operations relied on.

## SKILLS

---

**AI & Evaluation:** LLM evaluation, LLM-as-judge methodology, structured rubric design, rubric application, rubric calibration, RAG evaluation, eval pipeline design, prompt engineering

**Data & Analytics:** SQL, Python, documentation analytics, metrics strategy, content analytics, AI analytics

**Documentation & Strategy:** Documentation engineering, knowledge optimization, content governance, documentation quality, technical communication, knowledge management

**Domains:** AI/LLM evaluation infrastructure, knowledge engineering, AI agents, developer tools, HR technology, People Products

**Leadership & Process:** Cross-functional coordination, mentorship, project management, roadmap planning, public speaking

## EXPERIENCE

---

**Meta** — *Documentation Engineering Lead / Information Solutions Engineer*

*Nov 2021 - Jul 2026*

- **Hand-graded the verdict-balanced calibration set that anchored a production LLM judge** — personally scoring 200 prompts across five rubric dimensions (1,000 dimension-level judgments) against a four-verdict canonical vocabulary, applying the rubric consistently across the full batch to establish the ground truth the automated grader was measured against. The resulting judge reached  $\kappa = 0.86$  across 3,000 calibration calls.
- **Built and productionized a closed-loop AI evaluation system** that graded production AI agent outputs against a frozen five-dimension rubric and four-verdict scoring vocabulary, using three-replicate ensemble grading with per-dimension majority vote to hold judgments consistent across long batches. In a single quarter, the system graded 35,724 user prompts across 49 products, classified 914 knowledge gaps as hard or partial coverage failures, and generated 1,189 LLM-drafted remediations — corrected next steps routed through human review — at \$0.33 per gap, driving a **+9.3pp lift in eval pass rate** across supported products.
- **Lifted on-scope answer pass rate by +19.2pp** on one of the company's highest question-volume internal products — grading failing responses against the rubric to pinpoint the responsible content gaps, then verifying the lift with pre/post evaluation methodology now standard practice for the team. Cut content-staleness signal 79% in the same window.
- **Demonstrated that AI-agent answer quality is strongly predictable from retrieved-document characteristics alone** — analyzing a single research corpus of 200,000+ production prompt/response pairs to identify specific documentation traits as statistically significant predictors of positive user feedback, enabling targeted knowledge optimization rather than model tuning.
- **Pioneered AI-documentation metrics** — first in the company to measure AI consumption of docs at the per-document level, developing retrieval and citation metrics later adopted as company-wide standards. Revealed the inflection point where AI consumption surpassed human consumption, driving shifts in content and engagement strategy and shaping the function's two-year AI roadmap.

- **Designed and led a recurring company-wide training program** for 60+ documentation engineers on measurement, data storytelling, and responsible data use — building hands-on curriculum and coaching individuals and teams to standardize practices, strengthen privacy compliance, and communicate impact to leadership.
- **Diagnosed and resolved LLM failure modes in production** — redesigned prompting strategy, judge prompts, and citation-verification logic to cut false-positive abstentions in a retrieval-augmented generation evaluation pipeline.
- **Published a 10-skill Claude Code plugin suite** to the internal AI-coding-assistant marketplace supporting RAG evaluation, gap analysis, content health scoring, semantic clustering, golden-dataset generation, automated remediation workflows, report synthesis, and production AI quality measurement. Now standard tooling across Meta documentation teams.
- **Defined the documentation team's measurement standards, North Star goals, and OKRs** — leading metrics strategy and mentoring colleagues in SQL, analytics, and dashboarding. Promoted 3 times in 4 years (Junior → Staff).

**Living Computers: Museum + Labs** — *Staff Writer*

*Apr 2017 - Jul 2020*

- Managed and edited the museum's blog and refined its editorial voice with stakeholders; created technically accurate content with engineers and historians to make vintage computing systems legible to lay and expert audiences across written, video, and live formats.

**Geek Serv LLC** — *Founder / Owner / Operator*

*May 2008 - Apr 2017*

- Founded and scaled an IT services company serving thousands of clients across the Puget Sound region; built custom internal software with all accompanying documentation.

## EDUCATION

---

**B.A. Philosophy** — The Colorado College, Colorado Springs, CO